

An Introduction to Reinforcement Learning

Fundamental and formalism

Pei Liu

May 24, 2019

Department of Computer Science @UESTC

Table of contents

1. Background
2. Formalism
3. Algorithms
4. Extension

Background

Background - challenges

In supervised learning, the agent:

- learns from the labeled training set
- makes their outputs mimic the labels y given in the training set.

But it is indeed stuck in some scenarios, like *decision making* and *control problems*. Luckily, **Reinforcement Learning (RL)** came to the stage and solved them! In essence, it can be briefly summarized as:

- **experience-driven** autonomous learning
- improving over time through **trial and error**

Background - gallery

RL scenarios A: *Flappy Bird*




Figure 1: RL based Flappy Bird

Background - gallery

RL scenarios B: *Inverted Pendulum*




Figure 2: RL based balanced inverted pendulum

Background - question

In recent years, we have witnessed more examples like: AlphaGo, Self-driving, Robot, Control system.

Compared with previous generation RL, RL represented by above is more powerful in:

- directly recognizing *image* or *text* as the input
- more adaptive in the real world

It's called **Deep Reinforcement Learning** (DRL). DRL agents could be trained on raw, high-dimensional observations, solely based on a reward signal.

Background - question

Here our study of reinforcement learning will begin with a definition of the **Markov decision processes** (MDP), which provides the formalism in which RL problems are usually posed.

Formalism

Formalism - definition

A Markov decision process (MDP) is a tuple $(S, A, \{P_{sa}\}, \gamma, R)$, where:

- S, A are set of **states** and **actions**, respectively
- P_{sa} are the state transition probabilities
- $\gamma \in [0, 1)$ is called the **discount factor**
- $R: S \times A \mapsto \mathbb{R}$ is the **reward function**. (or $R: S \mapsto \mathbb{R}$)




Figure 3: The perception-action-learning loop

Formalism - procedure

At time t , the procedure in the loop of perception-action-learning would be as follows:

- the agent receives state S_t from the environment
- the agent uses its policy π to choose an action A_t
- once the action is executed, the environment transitions a step
- the environment provides the next state S_{t+1} and reward R_{t+1}




Figure 4: The perception-action-learning loop

Formalism - procedure of MDP

In MDP, the procedure could be simplified as:

- we starts in a state s_0
- we get to choose some action $a_0 \in A$ to take in the MDP
- once the action is executed, the state randomly transitions to some state $s_1 \sim P_{s_0 a_0}$
- we receive the reward $R(s_0)$ or $R(s_0, a_0)$
-

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$

Our total payoff is given by

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

Formalism - solution to MDP

In MDP, the procedure could be simplified as:

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$

Our goal in RL is to **choose actions** over time so as to **maximize the expected value** of the total payoff:

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

It is to say that we should find a **policy function** $\pi : S \mapsto A$ mapping from the states to the optimal actions. Under policy π , the total payoff we get is the **value function** as given by

$$V^\pi(s) = E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi]$$

Formalism - solution to MDP

With the definition of the **value function**

$$V^\pi(s) = E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi]$$

and a fixed **policy** function π . We can know that the value function $V^\pi(s)$ satisfies the **Bellman equations**:

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s')$$

which is also called as **Dynamic Programming**. Bellman's equations can be used to efficiently solve for V^π . (i.e. a set of $|S|$ linear equations in $|S|$ variables)

But the solution we want is the optimal value function $V^*(s)$.

Formalism - solution to MDP

We also define the **optimal value function** according to

$$V^*(s) = \max_{\pi} V^{\pi}(s). \quad (1)$$

In other words, this is the best possible expected sum of discounted rewards that can be attained using any policy. There is also a version of Bellman's equations for the optimal value function:

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s'). \quad (2)$$

The first term above is the immediate reward as before. The second term is the maximum over all actions a of the expected future sum of discounted rewards we'll get upon after action a . You should make sure you understand this equation and see why it makes sense.

We also define a policy $\pi^* : S \mapsto A$ as follows:

$$\pi^*(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s'). \quad (3)$$

Note that $\pi^*(s)$ gives the action a that attains the maximum in the “max” in Equation (2).

Algorithms

How can we find the optimal policy so as to maximize the expected total payoff?

Here we describe two efficient algorithms for solving **finite-state MDPs**:

- value iteration
- policy iteration

The two methods are heuristic and solved by iterations.

Algorithms - value iteration

In this case, the algorithm can be viewed as implementing a “Bellman backup operator” that takes a current estimate of the value function, and maps it to a new estimate.

1. For each state s , initialize $V(s) := 0$.

2. Repeat until convergence {

For every state, update $V(s) := R(s) + \max_{a \in A} \gamma \sum_{s'} P_{sa}(s') V(s')$.

}

This algorithm can be thought of as repeatedly trying to update the estimated value function using Bellman Equations (2).

Figure 5: Algorithm - value iteration

Algorithms - policy iteration

NOTE: step (a) can be done via solving Bellman's equations as described earlier, which in the case of a fixed policy, is just a set of $|S|$ linear equations in $|S|$ variables.

1. Initialize π randomly.
2. Repeat until convergence {
 - (a) Let $V := V^\pi$.
 - (b) For each state s , let $\pi(s) := \arg \max_{a \in A} \sum_{s'} P_{sa}(s')V(s')$.}

Figure 6: Algorithm - policy iteration

Extension

As we can see, MDP is a simple and ideal model. Even the transition probabilities P_{sa} is known!

Assuming that we don't know P_{sa} , and all we have is **just the reward signal** (actually, it's the situation), **how can we learn a agent/model from experiences via Reinforcement Learning?**

The **Inverted Pendulum** problem may be the best practice! See *here* for more details.

Next presentation, we will dive into **Flappy Bird** and reveal the principle **Q-learning** (a common method in RL) behind it!

Please be sure that you have known the solution of the problem **Inverted Pendulum**!

Questions?

References:

- Stanford CS229 - Reinforcement Learning
- cnblogs - An introduction to Reinforcement Learning
- A Brief Survey of Deep Reinforcement Learning